

# OTT REVENUE PREDICTION MODEL

A model to predict the revenue in million dollars based on the number of subscribers.

Data set:

Independent variable X: Subscribers/Year/Content Spend/Profit

Dependant variable Y: Overall revenue generated in dollars

```
In [1]: #import required libraries
import numpy as np
import pandas as pd
import os
import seaborn as sns
import matplotlib.pyplot as plt
```

```
In [2]: #List all files used
for dirname, _, filenames in os.walk(r"C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel"):
    for filename in filenames:
        print(os.path.join(dirname, filename))
```

```
C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\ContentSpend.csv
C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\NumSubscribers.csv
C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\NumSubscribersByRegion.csv
C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\Profit.csv
C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\Revenue.csv
C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\RevenueByRegion.csv
```

```
In [3]: #import dataset
df_profit = pd.read_csv(r"C:\Users\vdp10002\OneDrive - Advanced Micro Devices Inc\IISC_project\MainProject_OTTRe
df_subscribers = pd.read_csv(r"C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\NumSubscr
df_revenue = pd.read_csv(r"C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\Revenue.csv")
df_ContentSpend=pd.read_csv(r"C:\Users\vdp10002\Desktop\MainProject_OTTRevenuePredictionModel\Netflix\ContentSpe
#initial exploration of data
df_profit
```

Out[3]:

	Year	Profit
0	2012	0.050
1	2013	0.228
2	2014	0.403
3	2015	0.306
4	2016	0.379
5	2017	0.839
6	2018	1.600
7	2019	2.600
8	2020	4.500

In [4]: df\_subscribers

Out[4]:

	Year	Subscribers
0	2011	21.5
1	2012	25.7
2	2013	35.6
3	2014	47.9
4	2015	62.7
5	2016	79.9
6	2017	99.0
7	2018	124.3
8	2019	151.5
9	2020	192.9

In [5]: df\_revenue

Out[5]:

	Year	Revenue
0	2011	3.1
1	2012	3.5
2	2013	4.3
3	2014	5.4
4	2015	6.7
5	2016	8.8
6	2017	11.6
7	2018	15.7
8	2019	20.1
9	2020	24.9

In [6]: `df_ContentSpend`

Out[6]:

	Year	Content_spend
0	2016	6.88
1	2017	8.91
2	2018	12.00
3	2019	13.90
4	2020	11.80
5	2021	17.00

## VISUALIZE DATSET

```
In [7]: #understanding the trend of Revenue over the years, profit and content spend
x1 = df_profit['Year'].values
y1 = df_profit['Profit'].values

x2=df_subscribers['Year'].values
y2=df_subscribers['Subscribers'].values


x3=df_revenue['Year'].values
y3=df_revenue['Revenue'].values

x4=df_ContentSpend['Year'].values
y4=df_ContentSpend['Content_spend'].values

plt.rcParams["figure.figsize"] = (15,7)

plt.plot(x1, y1, 'red', label='Profit', marker='o', linestyle='-', linewidth='1')
#plt.plot(x2, y2, 'blue', label='Subscribers')
plt.plot(x3, y3, 'green', label='Revenue', marker='o', linestyle='-', linewidth='1')
plt.plot(x4, y4, 'black', label='Content Spend', marker='o', linestyle='-', linewidth='1')

plt.grid()
plt.xlabel('YEARS', fontweight="bold")
plt.ylabel('in Million Dollars', fontweight="bold")
plt.title('Revenue, Profit and Content Spend over the years', fontweight="bold")
plt.legend()
plt.show()
```



In [8]: *#understanding the trend of Revenue wrt # of Subscribers*

```
x5=df_subscribers['Year'].values
```

```
y5=df_subscribers['Subscribers'].values
```

```
plt.rcParams["figure.figsize"] = (15,7)
```

```
plt.plot(x5, y5, 'black', label='Subscribers', marker='o', linestyle='-', linewidth='1')
```

```
plt.grid()
```

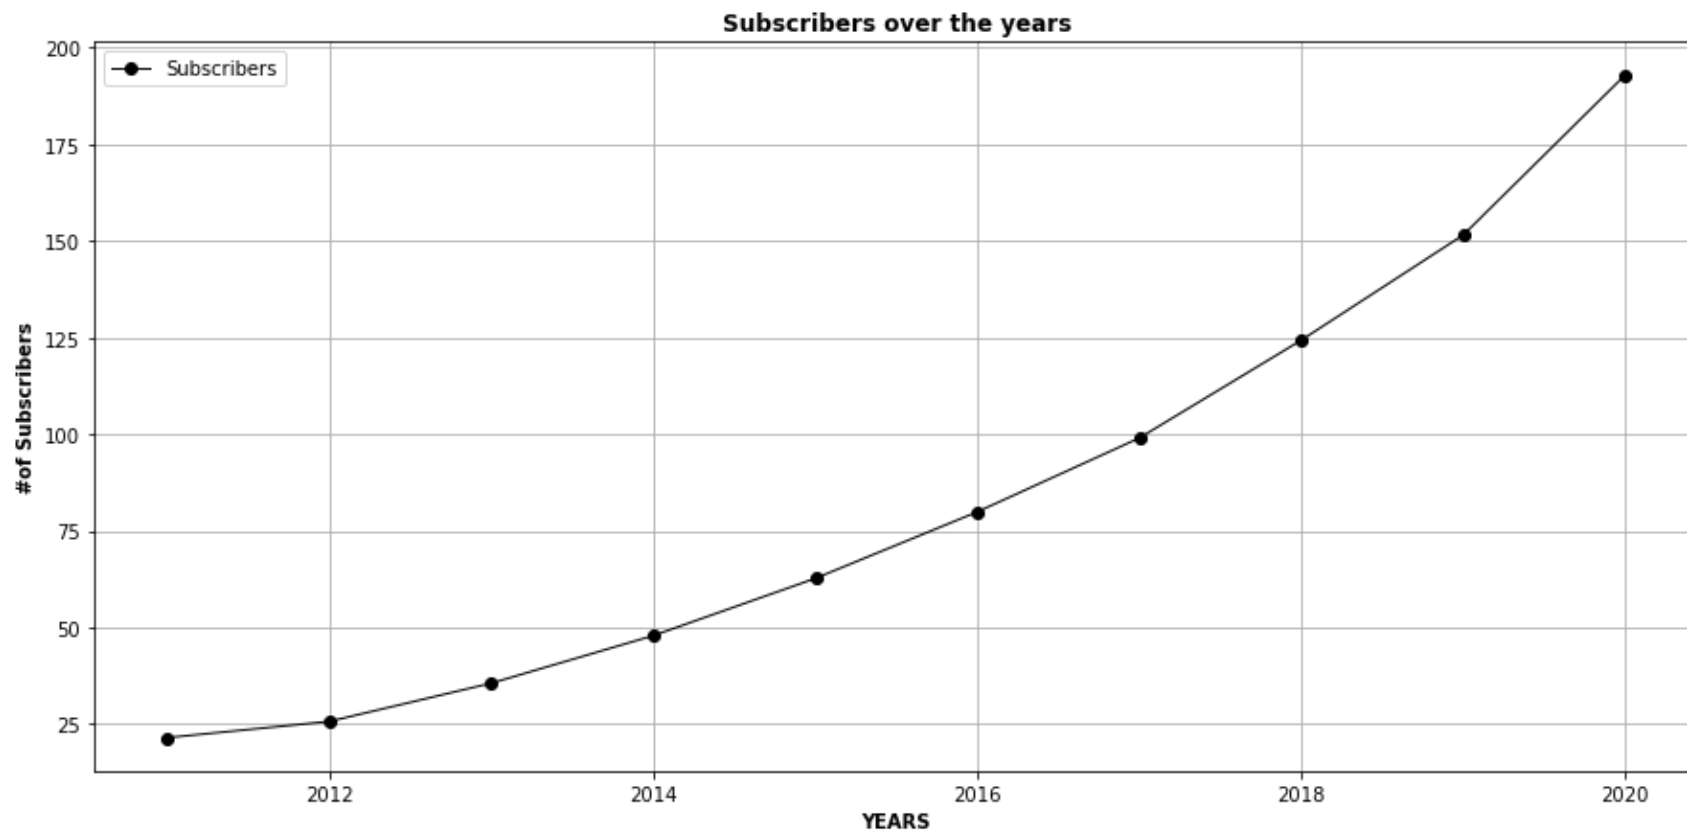
```
plt.xlabel('YEARS', fontweight="bold")
```

```
plt.ylabel('#of Subscribers', fontweight="bold")
```

```
plt.title('Subscribers over the years', fontweight="bold")
```

```
plt.legend()
```


```
plt.show()
```



```
In [9]: df_new = pd.merge(pd.merge(df_profit,df_revenue,on='Year'),df_subscribers,on='Year', how='right')
df_new1 = pd.merge(df_new, df_ContentSpend, on='Year', how='outer')
```

```
In [10]: # developing a histogram using DISPLOT
sns.displot(data = df_new1,
            x     = 'Revenue',
            height = 5,
            aspect = 2)
```


```
plt.show()
```






```
In [11]: sns.lmplot(x='Year', y='Revenue', data=df_new1)
```

```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x2d69d675f70>
```




```
In [12]: sns.lmplot(x='Subscribers', y='Revenue', data=df_new1)
```

```
Out[12]: <seaborn.axisgrid.FacetGrid at 0x2d69db47190>
```




```
In [13]: sns.lmplot(x='Content_spend', y='Revenue', data=df_new1)
```

```
Out[13]: <seaborn.axisgrid.FacetGrid at 0x2d69e488c10>
```



```
In [14]: sns.lmplot(x='Profit', y='Revenue', data=df_new1)
```

```
Out[14]: <seaborn.axisgrid.FacetGrid at 0x2d69e4faf10>
```



# MISSING VALUE ANALYSIS AND IMPUTATION

In [15]: `df_new1.isnull()`

Out[15]:

	Year	Profit	Revenue	Subscribers	Content_spend
0	False	True	True	False	True
1	False	False	False	False	True
2	False	False	False	False	True
3	False	False	False	False	True
4	False	False	False	False	True
5	False	False	False	False	False
6	False	False	False	False	False
7	False	False	False	False	False
8	False	False	False	False	False
9	False	False	False	False	False
10	False	True	True	True	False

In [16]: df\_new1

Out[16]:

	Year	Profit	Revenue	Subscribers	Content_spend
0	2011	NaN	NaN	21.5	NaN
1	2012	0.050	3.5	25.7	NaN
2	2013	0.228	4.3	35.6	NaN
3	2014	0.403	5.4	47.9	NaN
4	2015	0.306	6.7	62.7	NaN
5	2016	0.379	8.8	79.9	6.88
6	2017	0.839	11.6	99.0	8.91
7	2018	1.600	15.7	124.3	12.00
8	2019	2.600	20.1	151.5	13.90
9	2020	4.500	24.9	192.9	11.80
10	2021	NaN	NaN	NaN	17.00

```
In [17]: ## Replace all NaN values with 0
#df_new2= df_new1.fillna(0)
#df_new2
df_new2=df_new1
from sklearn.impute import SimpleImputer
imputer = SimpleImputer(missing_values=np.NaN, strategy='median')
print(imputer)
df_new2.Profit = imputer.fit_transform(df_new2['Profit'].values.reshape(-1,1))
df_new2.Revenue = imputer.fit_transform(df_new2['Revenue'].values.reshape(-1,1))
df_new2.Subscribers = imputer.fit_transform(df_new2['Subscribers'].values.reshape(-1,1))
df_new2.Content_spend = imputer.fit_transform(df_new2['Content_spend'].values.reshape(-1,1))
df_new2
```

SimpleImputer(strategy='median')

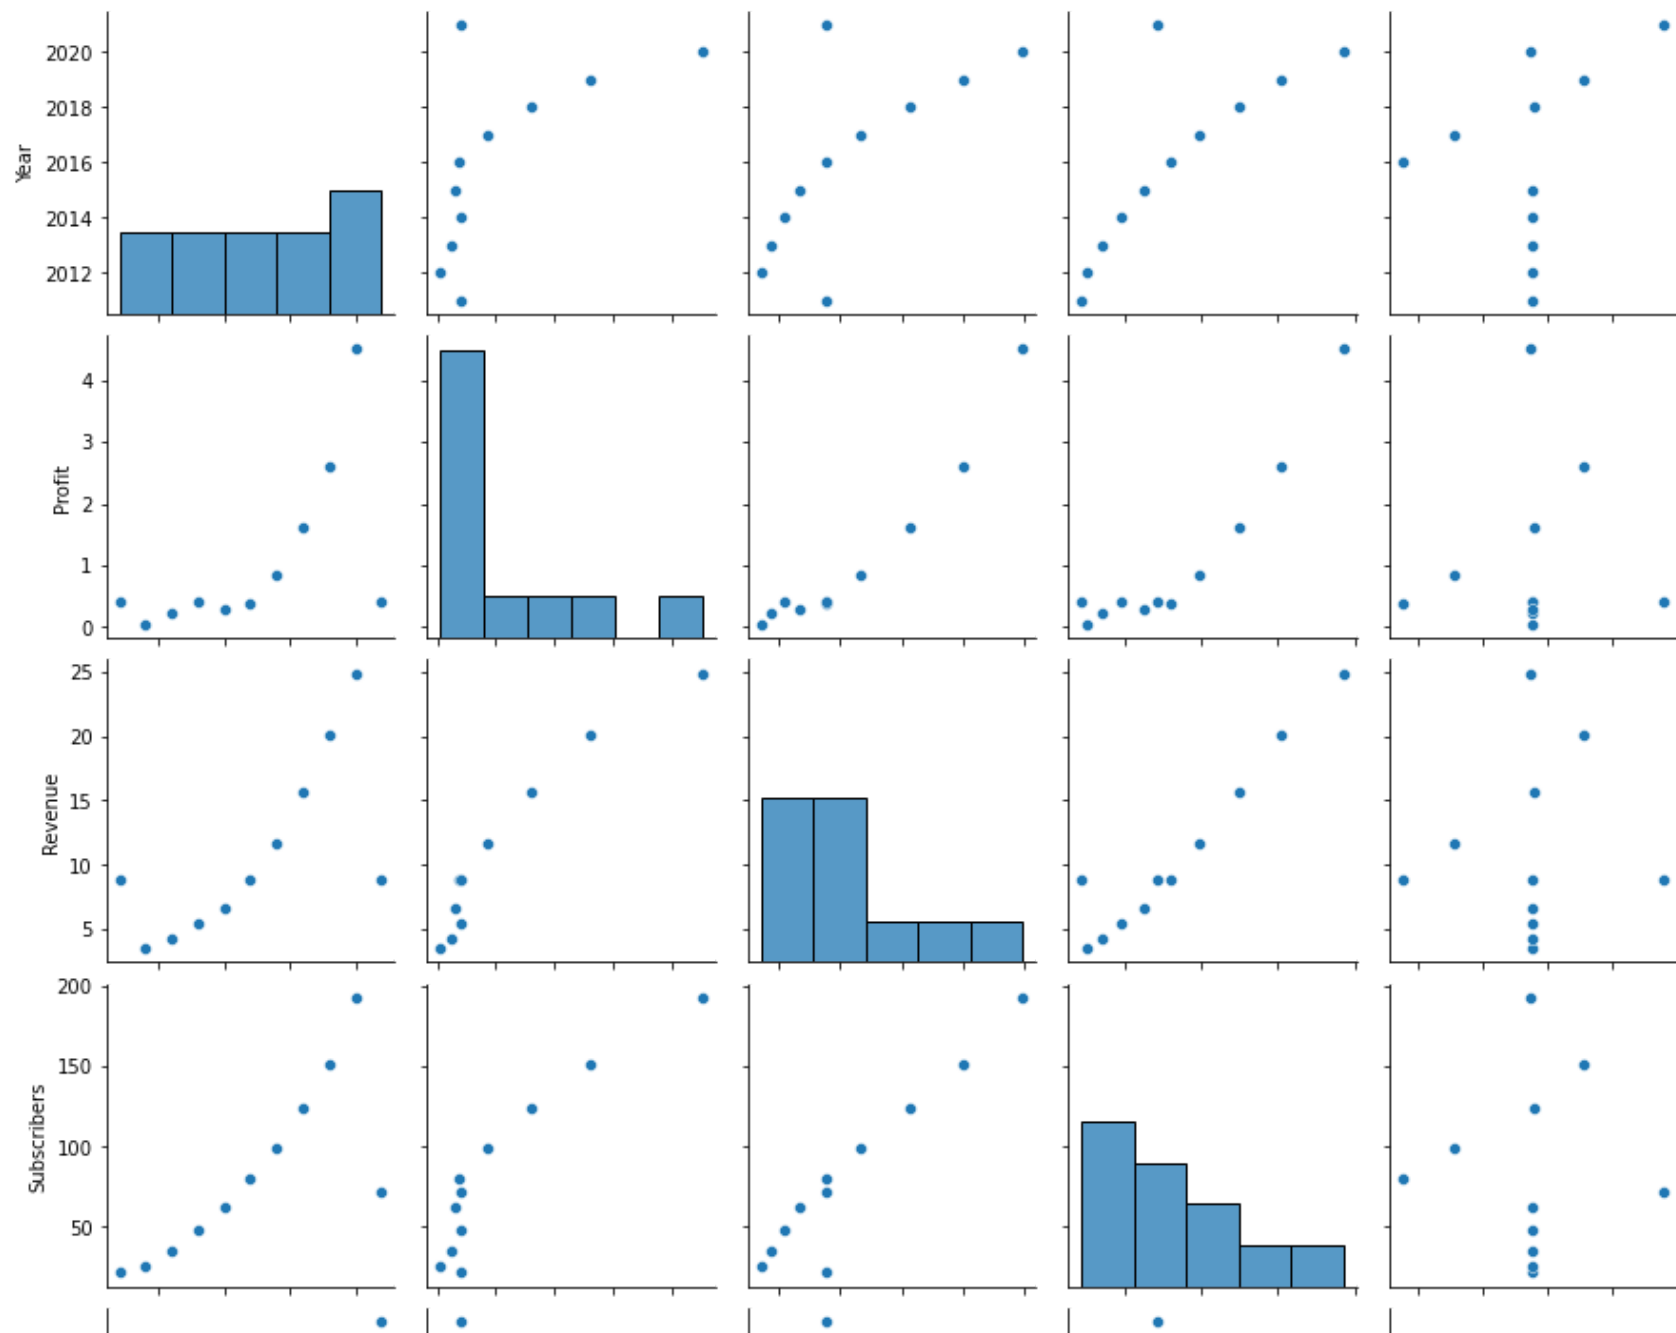
Out[17]:

	Year	Profit	Revenue	Subscribers	Content_spend
0	2011	0.403	8.8	21.5	11.90
1	2012	0.050	3.5	25.7	11.90
2	2013	0.228	4.3	35.6	11.90
3	2014	0.403	5.4	47.9	11.90
4	2015	0.306	6.7	62.7	11.90
5	2016	0.379	8.8	79.9	6.88
6	2017	0.839	11.6	99.0	8.91
7	2018	1.600	15.7	124.3	12.00
8	2019	2.600	20.1	151.5	13.90
9	2020	4.500	24.9	192.9	11.80
10	2021	0.403	8.8	71.3	17.00


## Understand the relationship between variables

```
In [18]: sns.pairplot(df_new2)
```

```
Out[18]: <seaborn.axisgrid.PairGrid at 0x2d69f927be0>
```







## OLS results interpretation

```
In [19]: from sklearn import linear_model
import statsmodels.api as sm

x = df_new2[['Year', 'Subscribers', 'Profit', 'Content_spend']]
y = df_new2['Revenue']

# adding a constant as an intercept is not included by default and has to be added manually
x = sm.add_constant(x)

model = sm.OLS(y, x).fit()

print_model = model.summary()
print(print_model)
```

### OLS Regression Results

```
=====
Dep. Variable:          Revenue    R-squared:                0.951
Model:                  OLS        Adj. R-squared:            0.919
Method:                 Least Squares    F-statistic:            29.37
Date:                  Tue, 29 Nov 2022    Prob (F-statistic):      0.000442
Time:                  22:31:32          Log-Likelihood:         -19.549
No. Observations:      11              AIC:                   49.10
Df Residuals:          6                BIC:                   51.09
Df Model:               4
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	513.1137	1205.304	0.426	0.685	-2436.158	3462.385
Year	-0.2550	0.602	-0.424	0.686	-1.727	1.217
Subscribers	0.0908	0.072	1.266	0.252	-0.085	0.266
Profit	1.7554	1.979	0.887	0.409	-3.087	6.598
Content_spend	0.1979	0.399	0.496	0.638	-0.779	1.174

```
=====
Omnibus:                4.600    Durbin-Watson:           1.356
Prob(Omnibus):           0.100    Jarque-Bera (JB):        1.808
Skew:                   0.959    Prob(JB):                 0.405
Kurtosis:               3.513    Cond. No.                 4.16e+06
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 4.16e+06. This might indicate that there are strong multicollinearity or other numerical problems.

C:\Users\vdp10002\Anaconda3\lib\site-packages\scipy\stats\stats.py:1541: UserWarning: kurtosistest only valid for n>=20 ... continuing anyway, n=11  
warnings.warn("kurtosistest only valid for n>=20 ... continuing ")

## Linear Regression Model

```
In [20]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn import preprocessing
from sklearn.metrics import r2_score
```

In [21]: *#split features: recognising dependent and independent variables*

```
y=df_new2[['Revenue']]
print(y)
x=df_new2.drop(['Revenue'], axis=1)
print(x)
```

	Revenue
0	8.8
1	3.5
2	4.3
3	5.4
4	6.7
5	8.8
6	11.6
7	15.7
8	20.1
9	24.9
10	8.8

	Year	Profit	Subscribers	Content_spend
0	2011	0.403	21.5	11.90
1	2012	0.050	25.7	11.90
2	2013	0.228	35.6	11.90
3	2014	0.403	47.9	11.90
4	2015	0.306	62.7	11.90
5	2016	0.379	79.9	6.88
6	2017	0.839	99.0	8.91
7	2018	1.600	124.3	12.00
8	2019	2.600	151.5	13.90
9	2020	4.500	192.9	11.80
10	2021	0.403	71.3	17.00

```
In [22]: #preparing training and test dataset
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.25)
x_train.shape
print(x_train)
```

	Year	Profit	Subscribers	Content_spend
2	2013	0.228	35.6	11.90
3	2014	0.403	47.9	11.90
9	2020	4.500	192.9	11.80
0	2011	0.403	21.5	11.90
6	2017	0.839	99.0	8.91
8	2019	2.600	151.5	13.90
5	2016	0.379	79.9	6.88
10	2021	0.403	71.3	17.00

```
In [23]: x_test
```

```
Out[23]:
```

	Year	Profit	Subscribers	Content_spend
4	2015	0.306	62.7	11.9
7	2018	1.600	124.3	12.0
1	2012	0.050	25.7	11.9

```
In [24]: y_test
```

```
Out[24]:
```

	Revenue
4	6.7
7	15.7
1	3.5

In [25]: y\_train

Out[25]:

	Revenue
2	4.3
3	5.4
9	24.9
0	8.8
6	11.6
8	20.1
5	8.8
10	8.8

In [26]: `%%time`  
*#instantiating linear regression model and fitting it to the training data*  
`LR = LinearRegression()`  
`LR.fit(x_train,y_train)`

CPU times: total: 0 ns  
Wall time: 2 ms

Out[26]: `LinearRegression()`

In [27]: `print('Intercept (c): ', LR.intercept_)`  
`print('Coefficient (m): ', LR.coef_)`  
*#scoring the model based on training and testing data*  
`LR_test_score=LR.score(x_test, y_test)`  
`LR_train_score=LR.score(x_train, y_train)`  
`print('LR Testing Score: ', LR_test_score)`  
`print('LR Trainig Score: ', LR_train_score)`

Intercept (c): [887.60609464]  
Coefficient (m): [[-0.44146596 1.38906515 0.10234273 0.29715788]]  
LR Testing Score: 0.9041527270851477  
LR Trainig Score: 0.9524730521774123

```
In [28]: #predicting on test data  
y_predict = LR.predict(x_test)  
print(y_predict)
```

```
[[ 8.43031193]  
 [15.23739231]  
 [ 5.61242811]]
```

```
In [29]: #evaluating Linear Regression model  
MSE_LR=mean_squared_error(y_test,y_predict)  
RMSE_LR=np.sqrt(mean_squared_error(y_test,y_predict))  
print("LR mean_sqrd_error is==", MSE_LR)  
print("LR root_mean_squared error of is==",RMSE_LR)
```

```
LR mean_sqrd_error is== 2.5567792534886373  
LR root_mean_squared error of is== 1.598993199950718
```

```
In [30]: #evaluating feature YEAR
x_yr=df_new2[['Year']]
y_yr=df_new2[['Revenue']]
x_yr_train, x_yr_test, y_yr_train, y_yr_test = train_test_split(x_yr, y_yr, test_size=0.25)
print('Shape of X_Year_TrainingData:', x_yr_train.shape)
LR_yr = LinearRegression()
LR_yr.fit(x_yr_train,y_yr_train)
y_yr_predict = LR_yr.predict(x_yr_test)
print(y_yr_predict)
print('Test Score:',LR_yr.score(x_yr_test, y_yr_test))
print('Train Score:', LR_yr.score(x_yr_train, y_yr_train))
print('Linear Model Coefficient (m): ', LR_yr.coef_)
print('Linear Model Coefficient (b): ', LR_yr.intercept_)
MSE_LR_yr=mean_squared_error(y_yr_test,y_yr_predict)
print("mean_sqrd_error is==", MSE_LR_yr)
```

```
Shape of X_Year_TrainingData: (8, 1)
[[12.29598145]
 [ 5.56615147]
 [13.41761978]]
Test Score: 0.5909239486505122
Train Score: 0.34682159254801936
Linear Model Coefficient (m): [[1.12163833]]
Linear Model Coefficient (b): [-2251.17017002]
mean_sqrd_error is== 20.170176505205415
```



In [31]: *#evaluating feature CONTENT SPEND*

```
x_cont=df_new2[['Content_spend']]
```

```
y_cont=df_new2[['Revenue']]
```

```
#print(x_sub)
```

```
#print(y_sub)
```

```
x_cont_train, x_cont_test, y_cont_train, y_cont_test = train_test_split(x_cont, y_cont, test_size=0.25)
```

```
#print('Shape of X_Subscribers_TrainingData:', x_sub_train.shape)
```

```
#print(x_sub_train)
```

```
LR_cont = LinearRegression()
```

```
LR_cont.fit(x_cont_train,y_cont_train)
```

```
y_cont_predict = LR_cont.predict(x_cont_test)
```

```
print(y_cont_predict)
```

```
print('Test Score:',LR_cont.score(x_cont_test, y_cont_test))
```

```
print('Train Score:', LR_cont.score(x_cont_train, y_cont_train))
```

```
print('Linear Model Coefficient (m): ', LR_cont.coef_)
```

```
print('Linear Model Coefficient (b): ', LR_cont.intercept_)
```

```
MSE_LR_cont=mean_squared_error(y_cont_test,y_cont_predict)
```

```
print("mean_sqrd_error is==", MSE_LR_cont)
```

```
[[11.33250145]
```

```
[12.01163664]
```

```
[10.89434971]]
```

```
Test Score: 0.014016633999728856
```

```
Train Score: 0.0037949537333521466
```

```
Linear Model Coefficient (m): [[0.21907587]]
```

```
Linear Model Coefficient (b): [8.28734688]
```

```
mean_sqrd_error is== 43.55636296159866
```

```
In [32]: #evaluating feature PROFIT
x_prof=df_new2[['Profit']]
y_prof=df_new2[['Revenue']]
#print(x_sub)
#print(y_sub)
x_prof_train, x_prof_test, y_prof_train, y_prof_test = train_test_split(x_prof, y_prof, test_size=0.25)
#print('Shape of X_Subscribers_TrainingData:', x_sub_train.shape)
#print(x_sub_train)
LR_prof = LinearRegression()
LR_prof.fit(x_prof_train,y_prof_train)
y_prof_predict = LR_prof.predict(x_prof_test)
print(y_prof_predict)
print('Test Score:',LR_prof.score(x_prof_test, y_prof_test))
print('Train Score:', LR_prof.score(x_prof_train, y_prof_train))
print('Linear Model Coefficient (m): ', LR_prof.coef_)
print('Linear Model Coefficient (b): ', LR_prof.intercept_)
MSE_LR_prof=mean_squared_error(y_prof_test,y_prof_predict)
print("mean_sqrd_error is==", MSE_LR_prof)

[[7.53681973]
 [7.65074992]
 [6.82000896]]
Test Score: 0.31356701329992187
Train Score: 0.9229748831552983
Linear Model Coefficient (m): [[4.74709117]]
Linear Model Coefficient (b): [5.73767217]
mean_sqrd_error is== 3.0889484401503524
```

```
In [33]: #evaluating feature SUBSCRIBER
x_sub=df_new2[['Subscribers']]
y_sub=df_new2[['Revenue']]
#print(x_sub)
#print(y_sub)
x_sub_train, x_sub_test, y_sub_train, y_sub_test = train_test_split(x_sub, y_sub, test_size=0.25)
#print('Shape of X_Subscribers_TrainingData:', x_sub_train.shape)
#print(x_sub_train)
LR_sub = LinearRegression()
LR_sub.fit(x_sub_train,y_sub_train)
y_sub_predict = LR_sub.predict(x_sub_test)
print(y_sub_predict)
print('Test Score:',LR_sub.score(x_sub_test, y_sub_test))
print('Train Score:', LR_sub.score(x_sub_train, y_sub_train))
print('Linear Model Coefficient (m): ', LR_sub.coef_)
print('Linear Model Coefficient (b): ', LR_sub.intercept_)
MSE_LR_sub=mean_squared_error(y_sub_test,y_sub_predict)
print("mean_sqrd_error is==", MSE_LR_sub)
```

```
[[7.39251431]
 [6.00928999]
 [9.05688178]]
```

```
Test Score: -3.31178195960747
```

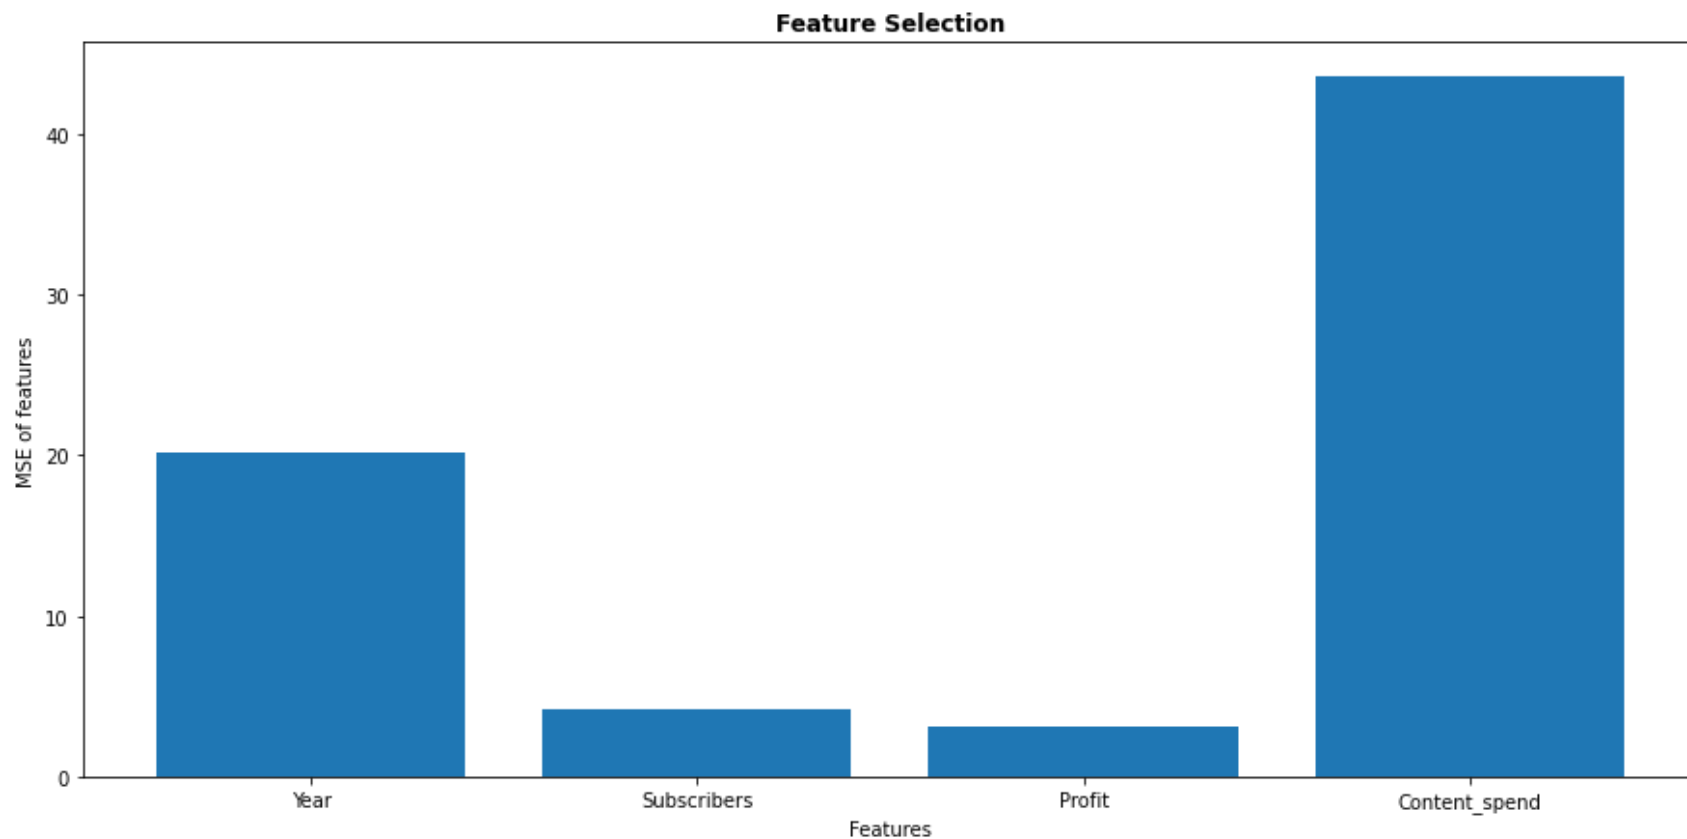
```
Train Score: 0.9057944001433338
```

```
Linear Model Coefficient (m): [[0.11245726]]
```

```
Linear Model Coefficient (b): [2.00581147]
```

```
mean_sqrd_error is== 4.148892418911189
```

```
In [34]: plt.bar(["Year", "Subscribers", "Profit", "Content_spend"], [MSE_LR_yr, MSE_LR_sub, MSE_LR_prof, MSE_LR_cont])  
plt.title("Feature Selection", fontweight='bold')  
plt.ylabel("MSE of features")  
plt.xlabel("Features")  
plt.show()
```



```
In [35]: # predict revenue with #ofSubscribers
y_predict_1 = LR_sub.predict([[200]])
y_predict_1
```

C:\Users\vdp10002\Anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(

```
Out[35]: array([[24.49726383]])
```

```
In [36]: # predict revenue for a particular year
y_predict_2 = LR_yr.predict([[2030]])
print(y_predict_2)
```

```
[[25.75564142]]
```

C:\Users\vdp10002\Anaconda3\lib\site-packages\sklearn\base.py:450: UserWarning: X does not have valid feature names, but LinearRegression was fitted with feature names  
warnings.warn(

## Random Forest Regression Model

```
In [37]: from sklearn.ensemble import RandomForestRegressor
```

```
In [38]: %%time
#instantiating Random Forest regression model and fitting it to the training data
RF = RandomForestRegressor(n_jobs=-1)
RF.fit(x_sub_train, y_sub_train)
```

```
CPU times: total: 109 ms
Wall time: 85 ms
```

<timed exec>:3: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples,), for example using ravel().

```
Out[38]: RandomForestRegressor(n_jobs=-1)
```

```
In [39]: # Calculate R2
RF_train_score=RF.score(x_sub_train, y_sub_train)
RF_test_score=RF.score(x_sub_test, y_sub_test)
print('RF Testing Score:',RF_test_score)
print('RF Taining Score:',RF_train_score)
```

```
RF Testing Score: -0.8282682448036791
RF Taining Score: 0.968001770168733
```

```
In [40]: #predicting on test data
y_predict_RF = RF.predict(x_sub_test)
print(y_predict_RF)
```

```
[5.514 5.514 8.647]
```

```
In [41]: #evaluating Random Forest regression model
score_RF=RF.score(y_sub_test,y_predict_RF)
MSE_RF=mean_squared_error(y_sub_test,y_predict_RF)
RMSE_RF=np.sqrt(mean_squared_error(y_sub_test,y_predict_RF))
print("RF r2 socre is ",score_RF)
print("RF mean_sqrd_error is==", MSE_RF)
print("RF root_mean_squared error of is==",RMSE_RF)
```

```
RF r2 socre is -0.6401499171377567
RF mean_sqrd_error is== 1.7592003333333206
RF root_mean_squared error of is== 1.3263484961854184
```

C:\Users\vdp10002\Anaconda3\lib\site-packages\sklearn\base.py:493: FutureWarning: The feature names should match those that were passed during fit. Starting version 1.2, an error will be raised.

Feature names unseen at fit time:

- Revenue

Feature names seen at fit time, yet now missing:

- Subscribers

```
warnings.warn(message, FutureWarning)
```

## KNN Regression Model

```
In [42]: from sklearn.neighbors import KNeighborsRegressor
```

```
In [43]: %%time
#instantiating KNN regression model and fitting it to the training data
KNN = KNeighborsRegressor(n_neighbors=2)
KNN.fit(x_sub_train,y_sub_train)
```

CPU times: total: 0 ns

Wall time: 2 ms

Out[43]: KNeighborsRegressor(n\_neighbors=2)

```
In [44]: KNN_test_score=KNN.score(x_sub_test, y_sub_test)
KNN_train_score=KNN.score(x_sub_train, y_sub_train)
print('KNN Testing Score:',KNN_test_score)
print('KNN Taining Score:',KNN_train_score)
```

KNN Testing Score: -1.9081986143187066

KNN Taining Score: 0.9104524248986294

```
In [45]: y_predict_knn = KNN.predict(x_sub_test)
print(y_predict_knn)
```

```
[[6.15]
 [6.15]
 [8.8 ]]
```

```
In [47]: score_KNN=KNN.score(y_sub_test,y_predict_RF)
MSE_KNN=mean_squared_error(y_sub_test,y_predict_knn)
RMSE_KNN=np.sqrt(mean_squared_error(y_sub_test,y_predict_knn))
print("r2 socre is ",score_KNN)
print("MSE is ",MSE_KNN)
print("RMSE is ",RMSE_KNN)
```

```
r2 socre is  -0.07644012559892577
MSE is  2.7983333333333342
RMSE is  1.6728219670166142
```

C:\Users\vdp10002\Anaconda3\lib\site-packages\sklearn\base.py:493: FutureWarning: The feature names should match those that were passed during fit. Starting version 1.2, an error will be raised.

Feature names unseen at fit time:

- Revenue

Feature names seen at fit time, yet now missing:

- Subscribers

```
warnings.warn(message, FutureWarning)
```

## Decision Tree Regression Model

```
In [48]: from sklearn.tree import DecisionTreeRegressor
```

```
In [49]: %%time
#instantiating Decision Tree regression model and fitting it to the training data
DT = DecisionTreeRegressor(random_state = 0)
DT.fit(x_sub_train,y_sub_train)
```

```
CPU times: total: 0 ns
```

```
Wall time: 4.04 ms
```

```
Out[49]: DecisionTreeRegressor(random_state=0)
```



```
In [50]: DT_test_score=DT.score(x_sub_test, y_sub_test)
DT_train_score=DT.score(x_sub_train, y_sub_train)
print('DT Testing Score:',DT_test_score)
print('DT Taining Score:',DT_train_score)
```

```
DT Testing Score: -2.0000000000000004
DT Taining Score: 1.0
```

```
In [51]: y_predict_DT = DT.predict(x_sub_test)
print(y_predict_DT)
```

```
[3.5 3.5 8.8]
```

```
In [52]: score_DT=DT.score(y_sub_test,y_predict_knn)
MSE_DT=mean_squared_error(y_sub_test,y_predict_DT)
RMSE_DT=np.sqrt(mean_squared_error(y_sub_test,y_predict_DT))
print("r2 socre is ",score_DT)
print("MSE is ",MSE_DT)
print("RMSE is ",RMSE_DT)
```

```
r2 socre is -1.9999999999999996
MSE is 2.8866666666666668
RMSE is 1.6990193249832881
```

C:\Users\vdp10002\Anaconda3\lib\site-packages\sklearn\base.py:493: FutureWarning: The feature names should match those that were passed during fit. Starting version 1.2, an error will be raised.

Feature names unseen at fit time:

- Revenue

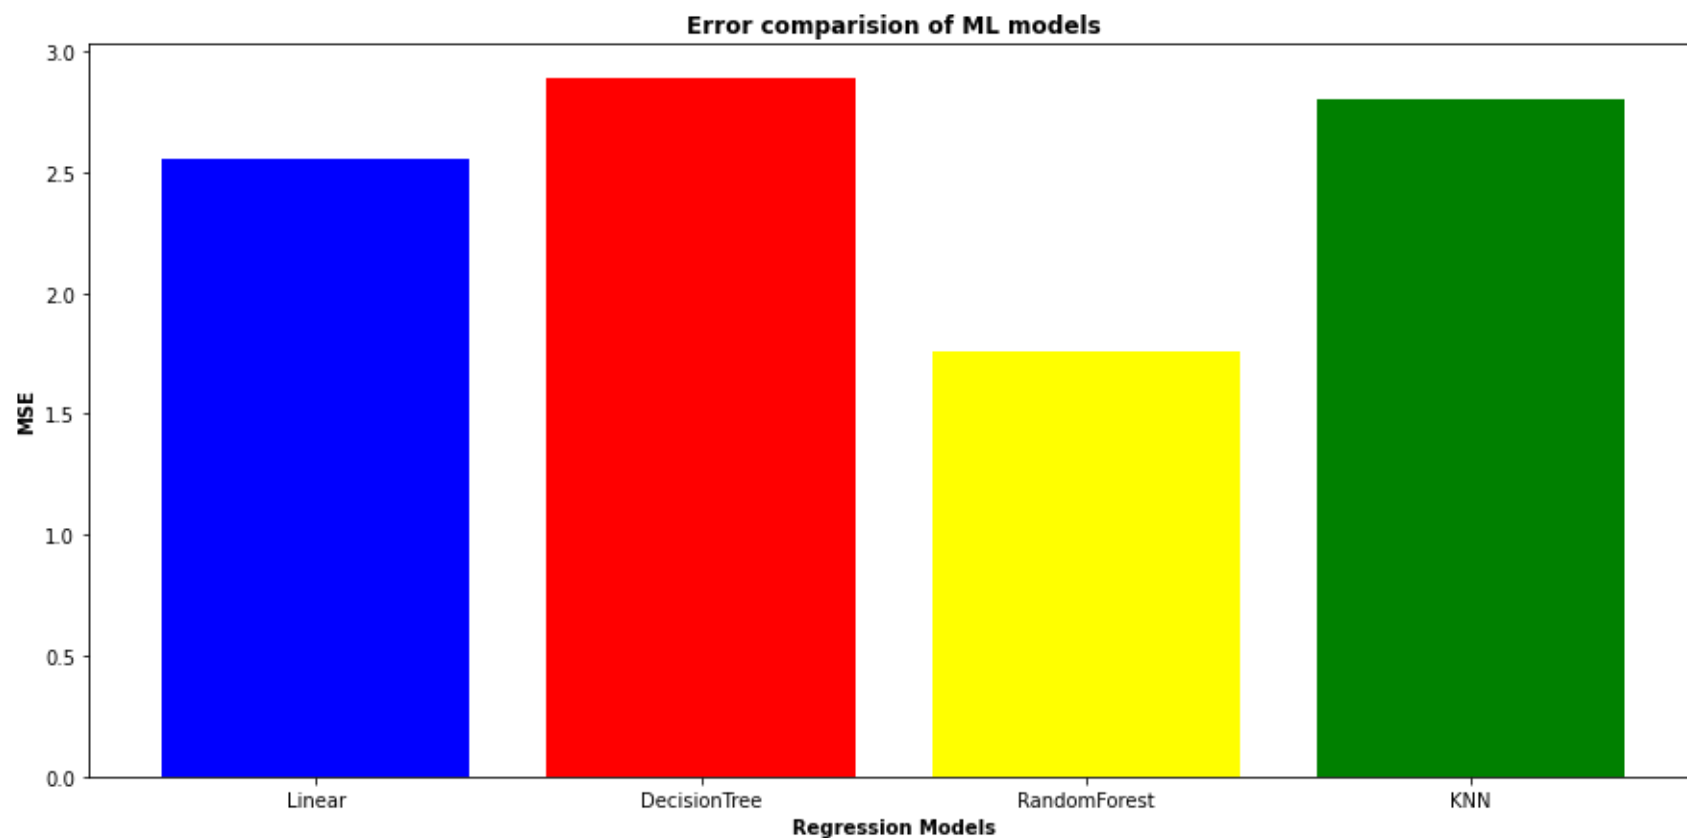
Feature names seen at fit time, yet now missing:

- Subscribers

```
warnings.warn(message, FutureWarning)
```

## Comparing Models

```
In [53]: plt.bar(["Linear", "DecisionTree", "RandomForest", "KNN"], [MSE_LR, MSE_DT, MSE_RF, MSE_KNN], color=['blue', 'red', 'yellow', 'green'])  
plt.title("Error comparision of ML models", fontweight="bold")  
plt.ylabel("MSE", fontweight="bold")  
plt.xlabel("Regression Models", fontweight="bold")  
plt.show()
```



In [ ]:

